

```

# Common EDA and plotting libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Allowing plots to appear in the notebook
%matplotlib inline

## Models
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier

## Model evaluators
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import precision_score, recall_score, f1_score

```

EDA

```

df = pd.read_csv("heart-disease.csv")
df.shape

```

```

(303, 14)

```

```

# Let's check the top 5 rows of our dataframe
df.head()

```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
0	63	1	3	145	233	1	0	150	0	2.3
1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

The top 10 rows

df.head(10)

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
0	63	1	3	145	233	1	0	150	0	2.3
1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6
5	57	1	0	140	192	0	1	148	0	0.4
6	56	0	1	140	294	0	0	153	0	1.3
7	44	1	1	120	263	0	1	173	0	0.0
8	52	1	2	172	199	1	1	162	0	0.5
9	57	1	2	150	168	0	1	174	0	1.6

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1
5	0	1	1
6	0	2	1
7	0	3	1
8	0	3	1
9	0	2	1

The bottom 10 Rows

df.tail(10)

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
293	67	1	2	152	212	0	0	150	0
294	44	1	0	120	169	0	1	144	1
295	63	1	0	140	187	0	0	144	1
296	63	0	0	124	197	0	1	136	1

297	59	1	0	164	176	1	0	90	0
1.0									
298	57	0	0	140	241	0	1	123	1
0.2									
299	45	1	3	110	264	0	1	132	0
1.2									
300	68	1	0	144	193	1	1	141	0
3.4									
301	57	1	0	130	131	0	1	115	1
1.2									
302	57	0	1	130	236	0	0	174	0
0.0									

	slope	ca	thal	target
293	1	0	3	0
294	0	0	1	0
295	2	2	3	0
296	1	0	2	0
297	1	2	1	0
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

Let's see how many positive (1) and negative (0) samples we have in our dataframe

```
df.target.value_counts()
```

```
1    165
```

```
0    138
```

```
Name: target, dtype: int64
```

Normalized value counts

```
df.target.value_counts(normalize=True)
```


```
1    0.544554
```

```
0    0.455446
```

```
Name: target, dtype: float64
```

Plot the value counts with a bar graph

```
df.target.value_counts().plot(kind="bar", color=["salmon", "lightblue"]);
```



Dataframe Summary

```
df.info()
```

```
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 303 entries, 0 to 302
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	age	303 non-null	int64
1	sex	303 non-null	int64
2	cp	303 non-null	int64
3	trestbps	303 non-null	int64
4	chol	303 non-null	int64
5	fbs	303 non-null	int64
6	restecg	303 non-null	int64
7	thalach	303 non-null	int64
8	exang	303 non-null	int64
9	oldpeak	303 non-null	float64
10	slope	303 non-null	int64
11	ca	303 non-null	int64
12	thal	303 non-null	int64
13	target	303 non-null	int64

```
dtypes: float64(1), int64(13)
```

```
memory usage: 33.3 KB
```

	age	sex	cp	trestbps	chol
fbf \					
count	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026
std	9.082101	0.466011	1.032052	17.538143	51.830751
min	29.000000	0.000000	0.000000	94.000000	126.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000
max	77.000000	1.000000	3.000000	200.000000	564.000000

	restecg	thalach	exang	oldpeak	slope
ca \					
count	303.000000	303.000000	303.000000	303.000000	303.000000
mean	0.528053	149.646865	0.326733	1.039604	1.399340
std	0.525860	22.905161	0.469794	1.161075	0.616226
min	0.000000	71.000000	0.000000	0.000000	0.000000
25%	0.000000	133.500000	0.000000	0.000000	1.000000
50%	1.000000	153.000000	0.000000	0.800000	1.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000

	thal	target
count	303.000000	303.000000
mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

Gender(Sex) variable counts
df.sex.value_counts()

```
1    207
0     96
Name: sex, dtype: int64
```

```
# Compare target column with sex column
pd.crosstab(df.target, df.sex)
```


```
sex      0      1
target
0        24    114
1        72     93
```

```
# Sex vs Target Plot
```

```
pd.crosstab(df.target, df.sex).plot(kind="bar", figsize=(10,6),
color=["salmon", "lightblue"])
```

```
# Plot attributes
```

```
plt.title("Heart Disease Frequency for Sex")
plt.xlabel("0 = No Disease, 1 = Disease")
plt.ylabel("Amount")
plt.legend(["Female", "Male"])
plt.xticks(rotation=0); # keep the labels on the x-axis vertical
```



```
# Age Vs Max Heart Rate Vs Target Plot
plt.figure(figsize=(10,6))
```

```
# Positive examples
```

```
plt.scatter(df.age[df.target==1],
            df.thalach[df.target==1],
```


```

c="salmon") # define it as a scatter figure

# Negative examples
plt.scatter(df.age[df.target==0],
            df.thalach[df.target==0],
            c="lightblue") # axis always come as (x, y)

# Plot attributes
plt.title("Heart Disease in function of Age and Max Heart Rate")
plt.xlabel("Age")
plt.legend(["Disease", "No Disease"])
plt.ylabel("Max Heart Rate");


```



```

# Age Histograms
df.age.plot.hist();


```



```
# Chest pain(cp) vs Target  
pd.crosstab(df.cp, df.target)
```

```
# Create a new crosstab and base plot  
pd.crosstab(df.cp, df.target).plot(kind="bar",  
                                   figsize=(10,6),  
                                   color=["lightblue", "salmon"])
```

```
# Plot attributes  
plt.title("Heart Disease Frequency Per Chest Pain Type")  
plt.xlabel("Chest Pain Type")  
plt.ylabel("Frequency")  
plt.legend(["No Disease", "Disease"])  
plt.xticks(rotation = 0);
```

Find the correlation between each pair variables

```
corr_matrix = df.corr()
corr_matrix
```

	age	sex	cp	trestbps	chol	
fbs \						
age	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308
sex	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032
cp	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444
trestbps	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531
chol	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294
fbs	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000
restecg	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189
thalach	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567
exang	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665
oldpeak	0.210013	0.096093	-0.149230	0.193216	0.053952	0.005747
slope	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894

ca	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979
thal	0.068001	0.210041	-0.161736	0.062210	0.098803	-0.032019
target	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046


	restecg	thalach	exang	oldpeak	slope	
ca \						
age	-0.116211	-0.398522	0.096801	0.210013	-0.168814	0.276326
sex	-0.058196	-0.044020	0.141664	0.096093	-0.030711	0.118261
cp	0.044421	0.295762	-0.394280	-0.149230	0.119717	-0.181053
trestbps	-0.114103	-0.046698	0.067616	0.193216	-0.121475	0.101389
chol	-0.151040	-0.009940	0.067023	0.053952	-0.004038	0.070511
fbs	-0.084189	-0.008567	0.025665	0.005747	-0.059894	0.137979
restecg	1.000000	0.044123	-0.070733	-0.058770	0.093045	-0.072042
thalach	0.044123	1.000000	-0.378812	-0.344187	0.386784	-0.213177
exang	-0.070733	-0.378812	1.000000	0.288223	-0.257748	0.115739
oldpeak	-0.058770	-0.344187	0.288223	1.000000	-0.577537	0.222682
slope	0.093045	0.386784	-0.257748	-0.577537	1.000000	-0.080155
ca	-0.072042	-0.213177	0.115739	0.222682	-0.080155	1.000000
thal	-0.011981	-0.096439	0.206754	0.210244	-0.104764	0.151832
target	0.137230	0.421741	-0.436757	-0.430696	0.345877	-0.391724

	thal	target
age	0.068001	-0.225439
sex	0.210041	-0.280937
cp	-0.161736	0.433798
trestbps	0.062210	-0.144931
chol	0.098803	-0.085239
fbs	-0.032019	-0.028046
restecg	-0.011981	0.137230
thalach	-0.096439	0.421741
exang	0.206754	-0.436757

```
oldpeak    0.210244 -0.430696
slope      -0.104764  0.345877
ca          0.151832 -0.391724
thal        1.000000 -0.344029
target     -0.344029  1.000000
```

Correlation Heatmap

```
corr_matrix = df.corr()
plt.figure(figsize=(15, 10))
sns.heatmap(corr_matrix,
            annot=True,
            linewidths=0.5,
            fmt= ".2f",
            cmap="YlGnBu");
```



```
df = pd.read_csv("heart-disease.csv") # 'DataFrame' shortened to 'df'
df.shape # (rows, columns)
```

```
(303, 14)
```

Modelling

Dataset processing

Everything except target variable

```
X = df.drop("target", axis=1)
```



```
C:\Users\zhizh\anaconda3\lib\site-packages\sklearn\linear_model\
_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>


Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
{'KNN': 0.6885245901639344,
 'Logistic Regression': 0.8852459016393442,
 'Random Forest': 0.8360655737704918}
```

Model Score Comparision Plot

```
model_compare = pd.DataFrame(model_scores, index=['accuracy'])
model_compare.T.plot.bar();
```



Logistic Regression have the best result

#Try improving the KNN model to produce better result

Create a list of train scores

```
train_scores = []
```

Create a list of test scores

```
test_scores = []
```

Create a list of different values for n_neighbors

```
neighbors = range(1, 21) # 1 to 20
```

Setup algorithm

```
knn = KNeighborsClassifier()
```

Loop through different neighbors values

```
for i in neighbors:
```

```
    knn.set_params(n_neighbors = i) # set neighbors value
```

```
    # Fit the algorithm
```

```

knn.fit(X_train, y_train)

# Update the training scores
train_scores.append(knn.score(X_train, y_train))

# Update the test scores
test_scores.append(knn.score(X_test, y_test))

train_scores


[1.0,
 0.8099173553719008,
 0.7727272727272727,
 0.743801652892562,
 0.7603305785123967,
 0.7520661157024794,
 0.743801652892562,
 0.7231404958677686,
 0.71900826446281,
 0.6942148760330579,
 0.7272727272727273,
 0.6983471074380165,
 0.6900826446280992,
 0.6942148760330579,
 0.6859504132231405,
 0.6735537190082644,
 0.6859504132231405,
 0.6652892561983471,
 0.6818181818181818,
 0.6694214876033058]

# Train Scores vs Test Scores Plot
plt.plot(neighbors, train_scores, label="Train score")
plt.plot(neighbors, test_scores, label="Test score")
plt.xticks(np.arange(1, 21, 1))
plt.xlabel("Number of neighbors")
plt.ylabel("Model score")
plt.legend()

print(f"Maximum KNN score on the test data: {max(test_scores)*100:.2f}%")

```

Maximum KNN score on the test data: 75.41%



```
# Tuning models with with RandomizedSearchCV
# Different LogisticRegression hyperparameters
log_reg_grid = {"C": np.logspace(-4, 4, 20),
                "solver": ["liblinear"]}

# Different RandomForestClassifier hyperparameters
rf_grid = {"n_estimators": np.arange(10, 1000, 50),
           "max_depth": [None, 3, 5, 10],
           "min_samples_split": np.arange(2, 20, 2),
           "min_samples_leaf": np.arange(1, 20, 2)}

# Setup random seed
np.random.seed(42)

# Setup random hyperparameter search for LogisticRegression
rs_log_reg = RandomizedSearchCV(LogisticRegression(),
                                param_distributions=log_reg_grid,
                                cv=5,
                                n_iter=20,
                                verbose=True)

# Fit random hyperparameter search model
rs_log_reg.fit(X_train, y_train);
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits


```

rs_log_reg.best_params_
{'solver': 'liblinear', 'C': 0.23357214690901212}
rs_log_reg.score(X_test, y_test)
0.8852459016393442

# Setup random seed
np.random.seed(42)

# Setup random hyperparameter search for RandomForestClassifier
rs_rf = RandomizedSearchCV(RandomForestClassifier(),
                           param_distributions=rf_grid,
                           cv=5,
                           n_iter=20,
                           verbose=True)

# Fit random hyperparameter search model
rs_rf.fit(X_train, y_train);

Fitting 5 folds for each of 20 candidates, totalling 100 fits

# Find the best parameters
rs_rf.best_params_
{'n_estimators': 210,
 'min_samples_split': 4,
 'min_samples_leaf': 19,
 'max_depth': 3}

# Evaluate the randomized search random forest model
rs_rf.score(X_test, y_test)
0.8688524590163934

# Different LogisticRegression hyperparameters
log_reg_grid = {"C": np.logspace(-4, 4, 20),
                "solver": ["liblinear"]}

# Setup grid hyperparameter search for LogisticRegression
gs_log_reg = GridSearchCV(LogisticRegression(),
                           param_grid=log_reg_grid,
                           cv=5,
                           verbose=True)

# Fit grid hyperparameter search model
gs_log_reg.fit(X_train, y_train);

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```

```

# Check the best parameters
gs_log_reg.best_params_

{'C': 0.23357214690901212, 'solver': 'liblinear'}

# Evaluate the model
gs_log_reg.score(X_test, y_test)

0.8852459016393442

# Make preidctions on test data
y_preds = gs_log_reg.predict(X_test)

y_preds
array([0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1,
       0,
       0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1,
       1,
       1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0],
      dtype=int64)

y_test
array([0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1,
       0,
       0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,
       1,
       1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0],
      dtype=int64)

from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt

# Predict the probabilities of the positive class
y_probs = gs_log_reg.predict_proba(X_test)[:, 1]


# Compute the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_probs)

# Compute the AUC
auc = roc_auc_score(y_test, y_probs)

# Plot the ROC curve
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, label='ROC curve (area = %.2f)' % auc)
plt.plot([-0.1, 1.1], [-0.1, 1.1], 'k--') # random predictions curve
plt.xlim([-0.1, 1.1])
plt.ylim([-0.1, 1.1])
plt.xlabel('False Positive Rate or (1 - Specificity)')
plt.ylabel('True Positive Rate or (Sensitivity)')
plt.title('Receiver Operating Characteristic')

```

```
plt.legend(loc="lower right")
plt.show()
```



```
# Display confusion matrix
print(confusion_matrix(y_test, y_preds))

[[25  4]
 [ 3 29]]

# Import Seaborn
import seaborn as sns
sns.set(font_scale=1.5) # Increase font size
```

```
def plot_conf_mat(y_test, y_preds):
    """
    Plots a confusion matrix using Seaborn's heatmap().
    """
    fig, ax = plt.subplots(figsize=(3, 3))
    ax = sns.heatmap(confusion_matrix(y_test, y_preds),
                      annot=True, # Annotate the boxes
                      cbar=False)
    plt.xlabel("true label")
    plt.ylabel("predicted label")

plot_conf_mat(y_test, y_preds)
```



```

                                scoring="accuracy") # accuracy as scoring
cv_acc
array([0.81967213, 0.90163934, 0.8852459 , 0.88333333, 0.75      ])
cv_acc = np.mean(cv_acc)
cv_acc
0.8479781420765027

# Cross-validated precision score
cv_precision = np.mean(cross_val_score(clf,
                                       X,
                                       y,
                                       cv=5, # 5-fold cross-validation
                                       scoring="precision")) #

precision as scoring
cv_precision
0.8215873015873015

# Cross-validated recall score
cv_recall = np.mean(cross_val_score(clf,
                                    X,
                                    y,
                                    cv=5, # 5-fold cross-validation
                                    scoring="recall")) # recall as


scoring
cv_recall
0.9272727272727274

# Cross-validated F1 score
cv_f1 = np.mean(cross_val_score(clf,
                                X,
                                y,
                                cv=5, # 5-fold cross-validation
                                scoring="f1")) # f1 as scoring

cv_f1
0.8705403543192143

# Visualizing cross-validated metrics
cv_metrics = pd.DataFrame({"Accuracy": cv_acc,
                           "Precision": cv_precision,
                           "Recall": cv_recall,
                           "F1": cv_f1},
                           index=[0])
cv_metrics.T.plot.bar(title="Cross-Validated Metrics", legend=False);

```



```
# Fit an instance of LogisticRegression (taken from above)
clf.fit(X_train, y_train);

# Check coef_
clf.coef_

array([[ 0.00369922, -0.90424089,  0.67472826, -0.0116134 , -
 0.00170364,
        0.04787688,  0.33490195,  0.02472938, -0.63120405, -
 0.57590942,
        0.47095136, -0.65165348, -0.69984206]])


# Match features to columns
features_dict = dict(zip(df.columns, list(clf.coef_[0])))
features_dict

{'age': 0.003699220351664148,
 'sex': -0.9042408930260735,
 'cp': 0.6747282624694215,
 'trestbps': -0.011613401789010375,
 'chol': -0.0017036441780094993,
```

```
'fbs': 0.047876883382302414,
'restecg': 0.3349019539205334,
'thalach': 0.024729383396378347,
'exang': -0.6312040510578483,
'oldpeak': -0.5759094230155162,
'slope': 0.47095135616471195,
'ca': -0.6516534832909596,
'thal': -0.6998420628111434}
```

```
# Visualize feature importance
```

```
features_df = pd.DataFrame(features_dict, index=[0])
features_df.T.plot.bar(title="Feature Importance", legend=False);
```



```
pd.crosstab(df["sex"], df["target"])
```

```
target    0    1
sex
0         24   72
1        114   93
```

```
# Contrast slope (positive coefficient) with target
```

```
pd.crosstab(df["slope"], df["target"])
```

target	0	1
slope		
0	12	9
1	91	49
2	35	107